

An Interactive Matchmaking Technique with Contextual Knowledge Elicitation

Xixi Xiao^{*}

Computer Science and Engineering Department
University of California, San Diego
x3xiao@eng.ucsd.edu

ABSTRACT

The term *Matchmaking* in computer science area refers to the process of using information technology to find the products or services that best fit with customer's individual preferences. For example, the modern computer dating system uses computer as a matchmaker to help match customers according to their tolerances and preferences. Online job hunting is another example for matchmaking system. Users input their preferences for job features and the search engine returns a set of candidate jobs. In this paper, we present an interactive model for matchmaking system, which uses contextual information to generate questions and gradually elicit useful information from customers. The information provided by customers in the interactive process in return helps improve matching accuracy.

Keywords

AsterixDB, query processor, question generation, semi-structured, Solrj, user interaction

1. INTRODUCTION

Matchmaking systems are widely used on the Internet to assist customers in finding fitting products. A typical matchmaking system is the job matching system. As the amount of information grows, traditional way of job searching is more and more costly. Recently, job seeking and recruiting websites have been experiencing a striking rise. Websites like *Monster.com*, *Indeed.com* are currently dominate this market.

One problem these systems are facing is the overload and inaccuracy of the output result set. One important reason for this deficiency is the insufficient information provided by the user. All users are asked to input the same kinds of information (e.g., location, position, etc.). No other attributes except these predefined ones could be used to express users' preferences. If you are looking for a "software engineer" position in the city "Mountain View" on *Monster.com*, it will

give you 1000+ results, which makes little sense.

Many researchers improve the matching performance by introducing data mining algorithms. Feature vectors are extracted and multiple similarity calculation methods has been proposed. To overcome the limitation of content-based text mining techniques, more and more people focus on digging out implicit semantic relationship between words to help improve matching accuracy.

However, in this paper, we proposed a novel idea to deal with the problem. We have designed and implemented an interactive matchmaking system, which utilizes contextual knowledge to elicit valuable information from users. The additional information provided by users can in return be used to pick up highly related products.

We have built an application in the scenario of job-hunting to illustrate our idea. A DSL (Domain Specific Language) has been designed as query language to talk to our matchmaking agent. User input a query, the agent finds a list of qualified jobs and figure out what attributes are important for the user to be competent for these jobs by analyzing related documents. Based on this context, the agent automatically generates a set of questions to elicit user to provide more useful and personalized information. With the user's answers, the agent further screens the jobs and create a personal profile for the user.

The datasets we conducted experiments on are a set of job documents and user profiles (semi-structured), which are manually acquired from available online resources. Different from general text documents, job description documents and user resumes have relative fixed structure, in other words, they are semi-structured data. Semi-structured data is a form of structured data that does not conform with the formal structure of data models associated with relational databases or other forms of data tables, but nonetheless contains tags or other markers to separate semantic elements and enforce hierarchies of records and fields within the data. Because of its flexibility, it's widely used in data representation and exchange system. Two typical types of semi-structured data are *XML* and *JSON*. The structure itself contains structural information that can be utilized for information indexing and retrieval.

AsterixDB, a result of about 3.5 years of R&D involving researchers at UC Irvine, UC Riverside, and UC San Diego, is a big data management system that supports ingesting, storing, managing, indexing, querying, and analyzing vast quantities of semi-structured information. With the help of *AsterixDB*, we are able to store and index hierarchical-structured data as well as extracting the structural informa-

^{*}A master student at University of California, San Diego

tion without knowing the exact schema of data. For example, if we have a job document contains a field "requirement", and the requirement has a field "minimum degree" with a corresponding value "Master", then we will extract the pair *Pair*("Jobs.requirement.minimum degree", "Master") from AsterixDB.

Another powerful platform we employed here is *Solr*, an open source enterprise search platform from the Apache Lucene project. *Solr* is very much like a key-value store. It is written in Java and runs as a standalone full-text search server within a servlet container such as Apache Tomcat or Jetty. Each document in *solr* is consisted of several field-value pairs. There is a configuration file in *solr*, which specifies the field names and types in the documents. We use *Solr* to index those path-value pairs extracted from AsterixDB. The combination of AsterixDB and *Solr* enables efficient query on both the text information and the structural information of the semi-structured data.

The rest of the paper is organized as follows. In section 2, we give a brief introduction to the related work. Section 3 describes the system architecture while section 4 details the design and implementation of each module. A concrete run-through example is given in section 5. Future work and possible improvements are states in section 6. Finally section 7 concludes the work.

2. RELATED WORK

Multiple job selection systems have been researched to improve the accuracy.

Some researches focus on the representation of knowledge in the given domain. In [6], an ontology-based method has been proposed for knowledge representation. The use of ontological descriptions for representing concepts according to a hierarchical structure helps exploit the semantic relationship between both job descriptions and user profiles. Based on the representation, a ranking algorithm allowing for qualification comparison has been designed. Maryam Fazel-Zarandi and Mark S. Fox presented an ontology-based hybrid approach in [2] to effectively match job seekers and job postings as well. A deductive model has been proposed and a similarity-based ranking algorithm has been designed for recommendation.

Personalized and dynamical recommendation system draws a lot of attentions as well. In [1], Ioannis Paparrizos, B. Barla Cambazoglu and Aristides Gionis trained a machine learning model by exploiting all past job transitions as well as the data associated with employees and institutions to predict an employee's next job transition. Wenxing Hong, Siting Zheng and Huan Wang [8] researched on a dynamic user profile-based job recommender system. Features are updated and extended dynamically based on the information and behaviors on the recruiting website which belong to the job applicants. Then a similarity calculation algorithm is proposed to match jobs with job applicants.

Many text mining techniques has been used for discovering more useful information from resumes(user profiles). For example, [5] classify the job applicants by using their resume in the method of text mining for improving the resource utilization and demand fulfill in the Resource Planning system that they build.

The recommendation algorithm is also the core of the system and vital to the accuracy. Basic approaches tend to be content-based. Features are extracted and similarity is

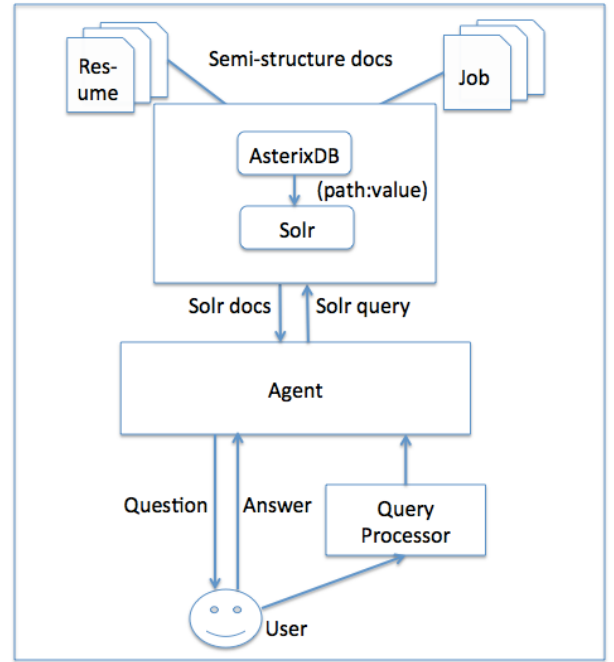


Figure 1: The system structure

computed, which can be used for ranking results. Collaborate filtering approach has been introduced to improve the performance. Hybrid approach has been proved to have a higher accuracy [4].

3. SYSTEM OVERVIEW

The architecture of the system is shown in Figure 1.

The semi-structured data, job data and resume data, are stored and indexed using *AsterixDB*. The JSON-like semi-structured data is actually a hierarchical data structure. It can be easily transformed to a tree structure. Each leaf node corresponds to a text field value while the root-to-leaf path is consisted of all the attributes, or tag names in *XML*, from the outer-most level to the inner-most level. We thus design an algorithm to extract all root-to-leaf paths and their corresponding text values from *AsterixDB*. In this way, the hierarchical data has been transformed to linear-structured data on which efficient indexing and searching methods can be performed. The path-value pairs are then organized and indexed in *Solr*. Each document in *solr* is consisted of all related path-value pairs and all theses pairs are indexed as key-value pairs.

A *DSL* has been designed for the query process. The syntax will be introduced in section 4. By parsing the user query, predicates(job constraints) can be acquired to form *solr* query sentences. Performing those queries in *solr* gives us a bunch of jobs the user may be interested in. If the number of jobs is within a reasonable range, the system can respond to the user query immediately. However, the number of jobs may either be too few, due to the strict constraints from user, or be too many if the user query is too general.

If the user query is too strict, *query relaxation* will be performed by the agent. The query relaxation techniques are usually used to expand query result set. Many algorithms have been designed to do query relaxation. For example, in [3], they focus on relaxing query conditions on numeric attributes. Say we are seeking for job-candidate pairs, such that the job and the candidate are in the same area (same zip code), the job's annual salary is at most 95K, and the candidate has at least 5 years of working experience. The result set for this query may be empty if no such pair of records satisfy both selection conditions. However, by relaxing both selection conditions on *salary* and *work year*, we can get a nonempty answer. For instance, for a condition *salary* ≤ 95, we could relax it to *salary* ≤ 100 or *salary* ≤ 120. In our work, the query relaxation is completed by dropping some predicates. Dropping predicates may lead to a large query result set, the same thing that will happen when the query conditions are too general. In this case, the interaction process will begin. During the interactive process, certain questions will be dynamically generated based on the current context and then delivered to the user. According to the answers provided by the user, a personal profile will be dynamically generated. User answers will be used to filter out non-related jobs as well since new predicates could be extracted. In a nutshell, the input of the system should be a user query that conforms to the syntax of the designed *DSL* while the output of the system are a personal profile and a reasonable number of jobs. The content of personal profile and the jobs being recommended are closely related to the user answers.

4. ALGORITHMS AND IMPLEMENTATION DETAILS

The system is mainly consisted of three modules: data module, query processor and interactive agent. The rest of this section will detail the implementation of each module.

4.1 Data module

As mentioned in Section 1, two existing platforms, *AsterixDB* and *solr* are used in this module to store and index data. We use *AsterixDB* to manage original semi-structured documents and we use *solr* to store and index the information extracted from *AsterixDB* for efficient information retrieval in the subsequent interactive process.

AsterixDB

There are three basic concepts-*dataverse*, *datatype* and *dataset* in *AsterixDB*. The top-level organizing concept in the *AsterixDB* world is the *dataverse*. A *dataverse*-short for "data universe"-is a place (similar to a database in a relational DBMS) in which to create and manage the types, datasets, functions, and other artifacts for a given *AsterixDB* application. A *datatype* tells *AsterixDB* what you know (or more accurately, what you want it to know) a priori about one of the kinds of data instances that you want *AsterixDB* to hold for you. A *dataset* is a collection of data instances of a *datatype*, and *AsterixDB* makes sure that the data instances that you put in it conform to its specified type. Since *AsterixDB* targets semistructured data, you can use open datatypes and tell it as little or as much as you wish about your data up front. Here is an example for the definition of datatypes *Job* and .

```
create type Job as open{
    id: int16,
    company_name: string,
    title: string,
    requirement: requirementType,
    location: string?
}

create type requirementType as open{
    minimum_degree: string?,
    discipline: [string]?
}
```

In this example, datatype *Job* and *requirementType* with several known fields is created. Some fields are of primitive types such like string, fields can also be of self-defining ADM(*AsterixDB* Data Model) types such like the *requirementType* here. There is a special *dataverse* in *AsterixDB* called *Metadata* in which there is a special dataset named *Dataset*. This dataset contains the *AsterixDB* system catalog. There is a record in this dataset for any dataset in any *dataverse*. For example, if we use a *dataverse searchAgent* to manage the job documents and user profiles, and we use two datasets-*Jobs* and *Resumes*-to store the corresponding data, then there will be a record in *Dataset* corresponds to the dataset *Jobs* in *dataverse searchAgent* and there will be another record in *Dataset* corresponds to dataset *Resumes* in *dataverse searchAgent*. Therefore, as long as we know the name of the *dataverse* used for managing our data, we are able to acquire the information of all datasets in that *dataverse*.

AsterixDB has provided us a *XQuery*-like query language to talk to it. With the information of datasets and the REST API provided by *AsterixDB*, we can get a JSON Object for each record in each dataset. By analyzing this JSON Object, several path-value pairs can be acquired, which we can use to index with *solr*. Say we have the following data record in the dataset that conforms to the *Job* datatype we defined above.

```
{
  "id":int16("101"),
  "company_name":"LinkedIn",
  "title":"Software Engineer",
  "require-
ment":{"minimum_degree":"Master","discipline":["CS"]},
  "location": "mountain view",
  "compensation": "$66/hour",
  "career level": "entry level"
}
```

Based on this record, we are able to extract several path-value pairs such as *Pair(dataset.requirement.discipline, "CS")*. I use *google-gson*, a JAVA library to help analyze the JSON Object got from the *AsterixDB* server. The algorithm for extracting information from *AsterixDB* and converting JSON Object to corresponding path-value pairs is described in Algorithm 1.

Algorithm 1: Extract path-value pairs

Require: Dataverse name D
Ensure: A set of path-value pairs P
1: Initialize $P \leftarrow \emptyset$,
2: Query Metadata of AsterixDB, $S \leftarrow$ get all datasets within the dataverse D .
3: **for** $d \in S$ **do**
4: $R \leftarrow$ All records(JSON Object) of dataset d
5: **for** each JSON Object $js \in R$ **do**
6: $P \leftarrow P \cup \text{DFS-JSON}(js, d.name)$
7: **end for**
8: **end for**
9: **return** P

The algorithm used in step 6 is very much like a DFS(depth-first-search) algorithm. A JSON object is actually a tree structured object, therefore by starting from the outer-most level(just like the root of the tree) and keeping track of the current path, which is actually a prefix of the final paths of all sub-objects, we can get all root-to-leaf paths by traversing the tree back and forth. The recursive algorithm is described in Algorithm 2.

Solr

Algorithm 2: DFS-JSON

Require: A JSON Object js , a string p denotes the current path prefix
Ensure: A set of path-value pairs P
1: Initialize $P \leftarrow \emptyset$
2: **for** $\{key, value\} \in js$ **do**
3: $p \leftarrow p + "." + key$
4: **if** $value$ is a JSONArray **then**
5: **for** $jse \in value$ **do**
6: **if** jse is a JSONObject **then**
7: DFS-JSON(jse, p)
8: **else**
9: $P \leftarrow P + \text{Pair}(p, value)$
10: **end if**
11: **end for**
12: **else if** $value$ is a JSONObject **then**
13: DFS-JSON($value, p$)
14: **else**
15: $P \leftarrow P + \text{Pair}(p, value)$
16: **end if**
17: $p \leftarrow$ drop the last part of p
18: **end for**
19: **return** P

Solr is a powerful and flexible standalone enterprise search server with a REST-like API. It supports many types of documents and various methods can be used to import data to Solr. We use *Solrj*, a java client, to programatically create documents to send to Solr.

The document sent to solr is very much like a key-value store, however in Solr, key is called *field*. There is a very important configuration file in Solr - the *schema.xml* file-which contains all of the details about which fields your documents can contain, and how those fields should be dealt with when adding documents to the index, or when querying those fields. This configuration file gives solr great flex-

ibility. A fantastic feature we employed when indexing the path-value pairs is the *Dynamic Fields* feature of solr. This feature enables on-the-fly addition of new fields. Since we extracts all path-value pairs and insert them into solr on-the-fly, this feature is exactly what we are looking for. For each job document, we construct a corresponding solr document and insert all related path-value pairs into it. We use the path as field name, which is stored and indexed, and the value is linked to the corresponding field. For user profiles, however, we use two fixed fields-path and word-to store the path-value pairs. The different designs are determined by the information retrieval demands in the subsequent interactive process. For job documents, we need to extract corresponding words of a given path, therefore we can use path as key to index text. Given a path value p , we can simply query solr with the condition "p:*" to retrieve all related documents. By iterating these documents and retrieve the values of field p , the text values corresponding to path p can be gathered. However, for user profiles, the interaction algorithm requires that given a specific text value, all corresponding paths can be retrieved. Some people may argue that we can inversely use text as indexed field. Unfortunately, the given text value may not be a single word, such like "software engineer", which cannot be indexed properly according to our experiments. An alternative way to index the path-value pairs is to use two fixed fields and use "AND" to join the two conditions. For example, we have a path-value pair $\text{Pair}(\text{Resumes.name}, \text{"Sissy"})$, the way we index it is to create a solr document {path: Resumes.name, word: "Sissy"}(there may be other fields in this documents). If we want to retrieve all corresponding path of the word "Sissy", we can use the query "word: Sissy" to retrieve all documents that contains the keyword "Sissy". By iterating these documents we can retrieve all values of the field path.

4.2 Query processor

The most ideal query and interactive language between human and computer is natural language. However, understanding natural language and extracting accurate information is challenging. For many dominant job hunting websites like *Monster*, *glassdoor*, a simple way to gather query information is to ask user to provide values for several fixed fields. For example, Monster.com asks user to fill in three text boxes-keywords, city and state-before starting matching jobs.

In this project, we have provided a more flexible query method. We have designed a DSL to facilitate the query process. It's more flexible and user friendly since the design of the syntax of DSL can vary from the actual requirements. We are able to design a DSL with the same grammar of natural language and we are able to inject logical operations between query conditions by special design. Here is the basic grammar for the user query language.

Query \rightarrow Seeking *Job*

Job \rightarrow Title *Company* Area Salary | Job or Job

Title \rightarrow TitleDictionary

Company $\rightarrow \epsilon$ | in *CompanyName* | in *CompanyType*
| preferably in *CompanyName* | preferably in *CompanyType*

$CompanyName \rightarrow StringConstant$
 $| CompanyName \text{ or } CompanyName$

$CompanyType \rightarrow CompanyTypeDictionary$
 $| CompanyType \text{ or } CompanyType$

$Area \rightarrow \epsilon \text{ in } AreaName \text{ around } AreaName$
 $| \text{preferably in } AreaName \text{ preferably around } AreaName$

$AreaName \rightarrow AreaDictionary$
 $| AreaName \text{ or } AreaName$

$Salary \rightarrow \epsilon \text{ with a salary of } Number \text{ per } Unit$

$StringConstant \rightarrow (["a" - "z", "A" - "Z", "_", "-", "0" - "9"])^+$

$Number \rightarrow (["0" - "9"])^+$

$Unit \rightarrow "month" \text{ } | \text{ } "hour" \text{ } | \text{ } "year"$

We omitted the definition for *TitleDictionary*, *CompanyTypeDictionary* and *AreaDictionary* here, which are lists of candidate words. The above grammar helps to check whether an input user query has the correct syntax or not. With the help of *JavaCC*, we have implemented a parser for this DSL. When a valid user query comes, the parser will generate an AST (Abstract Syntax Tree). By navigating along the AST, query conditions and logical operations can be extracted, which can be then transformed to solr query sentences.

For example, here is a valid user query:

"seeking a software engineer position preferably around San Francisco or Mountain View or San Diego"

By parsing this query, several path-value pairs can be extracted:

Pair("Jobs.title", "software engineer")
 Pair("Jobs.location", "San Francisco")
 Pair("Jobs.location", "Mountain View")
 Pair("Jobs.location", "San Diego")

Besides these pairs, the keywords "preferably", "around" will be recorded. The corresponding Solr query will be:

Jobs.title: "software engineer" AND (Jobs.location:"San Francisco" OR Jobs.location:"Mountain View" OR Jobs.location:"San Diego")

The keywords "preferably" and "around" here can be utilized for query relaxation.

4.3 Interactive agent

A collection of qualified job documents will be retrieved according to the input user query. As we discussed before, a large result set makes little sense. By interacting with user

may provide us more useful information and help us figure what exactly the user is looking for.

The interaction process involves asking user questions and collecting their answers. The questions should be generated dynamically and the user answers should be used to narrow qualified job documents set. Meanwhile, a personal profile for the user will be automatically generated based on user's answers to those corresponding questions.

Question generation

Before generating questions, we need to figure out what attributes are essential for the user to get the job he or she desires. The algorithm we designed to get those attributes is described in Algorithm 3.

Algorithm 3: Extract attributes

Require: A collection of job documents J , a frequency threshold θ

Ensure: A set of paths and their corresponding histograms $R = \{p, h\}$

- 1: Initialize $P \leftarrow \emptyset, R \leftarrow \emptyset, T \leftarrow \emptyset$
- 2: **for** $d \in J$ **do**
- 3: $T.wordlist \leftarrow T.wordlist +$ all field values of d
- 4: Update the number of occurrences of all field values of d
- 5: **end for**
- 6: According to T , construct a histogram H to record the number of occurrences of each word/phrase. All related paths to a specific word/phrase should be recorded as well.
- 7: Sort the words/phrases by their occurrence frequencies
- 8: $W \leftarrow \emptyset$
- 9: **for** $u \in H.wordlist$ **do**
- 10: **if** $u.frequency > \theta$ **then**
- 11: $W \leftarrow W + u$
- 12: **end if**
- 13: **end for**
- 14: $P \leftarrow \emptyset$
- 15: **for** $w \leftarrow$ each word in W **do**
- 16: $P \leftarrow P +$ all paths related to w in any user profiles indexed in Solr
- 17: **end for**
- 18: **for** $p \in P$ **do**
- 19: Construct a histogram h to record the number of occurrences of all values related to p in any user profiles, $R \leftarrow R + \{p, h\}$
- 20: **end for**
- 21: **return** R

The basic idea of the algorithm is that if a path in a user profile corresponds a word/phrase that appears several times in the job documents the user may be interested in, then the path, or we can say the attribute, may be essential. For example, the word "JAVA" appears many times in several qualified job documents, and we find that people usually mention the word "JAVA" when they stating their "programming language", then it's reasonable to say that "programming language" is an important attribute for job matching. Therefore, a proper question to ask the user is "What pro-

gramming language are you familiar with?”.

Job refilter

User answers provide us with additional information, which contributes to more accurate job matching. As we described above, a question is generated based on a path and for each path, we have constructed a histogram to record expected values for that path. If the user answer matches one of those high frequency words/phrases, we'll add this information to user's profile and use this answer for filtering jobs. Otherwise, questions will be designed to further explore the user's abilities. Algorithm 4 describes how we utilize user answers to form more query predicates and refilter candidate jobs.

Algorithm 4: Job refilter

Require: A set of words W , a collection of current qualified jobs J , a set of words and all paths related to each word $S = \{w, P\}$
Ensure: A collection of jobs R

- 1: Initialize $R \leftarrow \emptyset, C \leftarrow \emptyset$
- 2: **for** $u \in W$ **do**
- 3: **for** $\{w, P\} \in S$ **do**
- 4: **if** $u = w$ **then**
- 5: **for** $p \in P$ **do**
- 6: $C \leftarrow C + \text{Pair}(p, w)$
- 7: **end for**
- 8: **end if**
- 9: **end for**
- 10: **end for**
- 11: Form solr query according to C and J , get new job collection R
- 12: **return** R

Profile generation

A set of question-answer pairs, in other words, path-value pairs can be acquired during the interactive process. These pairs record the information that are essential for the user to be qualified for the jobs he or she is looking for. Therefore, we automatically organized and formatted the information to create a personal profile for the user, which is similar to a resume.

The interactive agent is just like a state machine. It maintains the current context, generate questions based on the current context to interact with user and then update the context according to user's feedback.

5. RUN-THROUGH EXAMPLE

In this section, we are going to run through a concrete example step by step to show how the application works. We argue that the interactive process returns a better fitted job sets than merely using the original input user query to select jobs.

• Step 1: Input user query

User should input a valid query which conforms to the grammar we designed for the DSL. We choose the following query:

”seeking a software engineer position around San Francisco or Mountain View or San Diego”

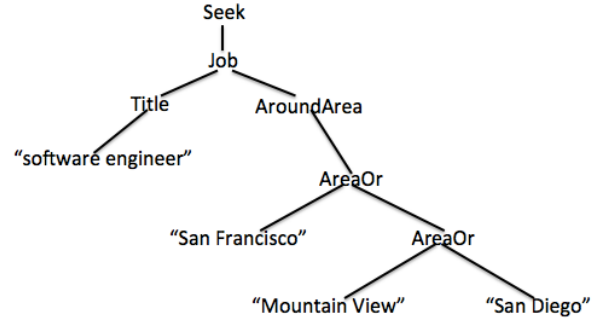


Figure 2: AST

The query is parsed to an abstract syntax tree as shown in Figure 2.

Each internal node is associated with a node type, which implies how to proceed while each leaf node is associated with a text value. For example, when we encounter the internal node of the type *AroundArea*, we will record the keyword "around". This keyword implies that if we can not find matching jobs exactly at the given area, we could expand the area to some extent. When we encounter the node type *Title*, we would know that the child of this node is a leaf node which stores the title information. Since we know beforehand that there would be a path "Jobs.title" in every job documents, we would store the pair *Pair*(Jobs.title, "software engineer").

After traversing the AST of the given query, we get the following set of path-value pairs:

```

Pair("Jobs.title", "software engineer")
Pair("Jobs.location", "San Francisco")
Pair("Jobs.location", "Mountain View")
Pair("Jobs.location", "San Diego")
  
```

• Step 2: Extract a set of jobs the user may be interested in

The above path-value pairs, which are actually query predicates, then are analyzed and combined to a single Solr query sentence:

```

Jobs.title: "software engineer" AND (Jobs.location:"San Francisco" OR Jobs.location:"Mountain View" OR Jobs.location:"San Diego")
  
```

By querying Solr, a set of qualified jobs will be returned.

• Step 3: Analyze current qualified jobs as well as the resume corpus, figure out related attributes

Following algorithm 3, firstly a histogram of vocabularies appearing in any field of any job documents in the current qualified job set is constructed. Each vocabulary(a word or a phrase) has a frequency, that is, the number of occurrences




Figure 3: Histogram of vocabularies from job documents

of this word/phrase divided by the total number of occurrences of all vocabularies. Figure 3 shows the histogram (we omitted the low-frequency vocabularies here).

Vocabulary with a frequency greater than the threshold (a value between 0 1, which we set up at the very beginning) is selected. Therefore, we get a list of words/phrases that appear multiple times in the qualified jobs. Then we use these words/phrases to figure out what attributes are essential. As stated in step 16 of algorithm 3, for each word/phrase, we retrieve all related paths in any resumes in our data corpus. As we discussed in section 4, it's reasonable to assert that the attributes corresponding to these paths are important.

• **Step 4: Generate questions based on the paths list** After we get the paths list, following step 19 of the algorithm, we construct a histogram of each path, recording the number of occurrences of all vocabularies related to it. 4 shows the distribution of vocabularies related to the path "resumes.experience.title". These words/phrases are the expected values of the corresponding path.

The paths are sorted according to their frequencies. We start with the top one—"resumes.education.degree". The question prompted to the user is:

"Could you tell me more about your education experience, what degree did you get?"

User gives an answer to the question, say "Master". The conversation continues with the next question.

"Could you tell me more about your working experience, what was your title?"

We have demonstrated the histogram for the path "resumes.experience.title" in Figure 4. There are a list of expected vocabularies. If the user answer matches any one




Figure 4: Histogram of vocabularies related to a specific path

of them, the next path will be chosen and corresponding question will be generated. However, if the user's answer doesn't match any expected value, the agent starts to elicit user to give more accurate information. For example, in this case, if the user answers "Design Technician", which is not among the vocabulary list, the agent picks the top vocabulary and generate a question in the pattern of "Do you have ..." or "Have you ever been a ...", etc. Here the question is:

"Have you ever been a software engineer?"

If the user's answer is "yes", the agent record the answer and proceed to the next path. Otherwise, a negative effect is recorded. When the negative effect reaches a pre-determined value, the agent will urge the user to try other kinds of job. Suppose the user gives a positive answer here, the conversation continues:

Agent: "Could you tell me more about your location, what city are you living in?"

User: "San Diego"

Agent: "Could you tell me more about your education experience, what is the major?"

User: "Computer Science"

Agent: "What skills do you have?"

User: "Java, C++"

The conversation ends here. A personal profile is automatically generated for the user based on their answers. • **Step 5: Use user answers to further screen jobs**

6. DISCUSSION

The idea of this interactive search agent is novel and the implementation is a big challenge. The system actually involves many computer science topics.

The first part relates to the representation of the knowledge. As we stated in section 2, several methods have been proposed including ontological representation. The advantage

of ontology-based knowledge representation is obvious. The hierarchical structure itself contains semantic information. Machine learning algorithms can be easily applied to it. It is possible to integrate the idea with our system and apply related machine algorithm to pre-process attribute values. There are actually a lot of text-processing related work we can do to improve the matching accuracy. We are now doing the exact word matching. However, the existence of synonyms asks for clustering words that have similar meaning. For example, "software engineer" and "programmer" should be considered in the same group. Many researchers in the data mining field have dealt with this issue.

Besides, the DSL provides great flexibility for the searching and matching process. For example, the keyword "preferably" can be used in multiple ways. One way I could think about is to assign a weight to the corresponding query conditions. The weight information tells the agent which features are more important and which are less, which we believe can be utilized to design suitable algorithms to improve performance.

A big problem we faced during the design process is how to generate proper easy-understanding questions. [9], [7] proposed some ideas to deal with the question generation problem. Unfortunately, none of those works address our problem. Firstly, in our system, we need to generate a question given a specific path. It's much harder than transforming a narrative sentence to a corresponding question. Different sentence patterns should be applied to different vocabularies. Besides, the order of the question should be concerned as well. For example, we may have two paths, "Education.university_name" and "Education.graduation_date", based on which we need to generate questions. In the correct logic, the question should be "Could you tell me more about your education, what university have you attended? When did you graduate from the university?" rather than "Could you tell me more about your education, when did you graduate from the university? What university have you attended?". This is a crossing area of linguistic and computer science, which makes it really challenging.

7. CONCLUSION

In this paper, we presents an interactive job searching and recommendation system.

Combined with the content-based information, we utilized the structural information of user profiles to improve matching accuracy. We employed two powerful databases-AsterixDB and Solr-to store and index data. We designed the algorithm to retrieve structural information from AsterixDB as well as configuring the proper indexing schema for Solr according to the searching demands.

We have proposed an interactive model to improve matching accuracy. The interactive process between user and agent helps discover important attributes and prob implicit features of the user. This is actually a bilateral selection process. Users express their preferences for the job, by analyzing job data and user constraints, important vocabularies are extracted, based on we are able to figure out essential attributes for the user to get the desired job. The user answers in return help the agent find better fitted jobs for them.

The automatically generated personal profile is a bonus of this system. The entries in the profiles are exactly those essential attributes the user should have to be competent for

the jobs they desired.

8. REFERENCES

- [1] Ioannis Paparrizos, B. Barla Cambazoglu, Aristides Gionis. Machine Learned Job Recommendation. *RecSys*, 2011.
- [2] Maryam Fazel-Zarandi, Mark S. Fox. Semantic Matchmaking for Job Recruitment: An Ontology-Based Hybrid Approach. In *Proceedings of the 3rd International Workshop on Service Matchmaking and Resource Retrieval*, 2009.
- [3] Nick Koudas, Chen Li, Anthony K. H. Tung, Rares Vernica. Relaxing Join and Selection Queries. *VLDB '06 Proceedings of the 32nd international conference on Very large data bases*, pages 199–210, 2006.
- [4] Prem Melville, Raymond J. Mooney and Ramadass Nagarajan. Content-Boosted Collaborative Filtering. In *Proceedings of the Eighteenth National Conference on Artificial Intelligence*, pages 187–192, 2002.
- [5] Tere Gonzalez, Pano Santos, Fernando Orozco, Mildreth Alcaraz, Victor Zaldivar, Alberto De Obeso, Alan Garcia. Adaptive Employee Profile Classification for Resource Planning Tool. In *Proceedings of SRII Global Conference 2012 Annual*, pages 544–553, 2012.
- [6] Valentina Gatteschi, Fabrizio Lamberti, Andrea Sanna, Claudio Demartini. A Semantic Matchmaking System For Job Recruitment. *Proceedings of I-KNOW 2010*, pages 50–59, 2010.
- [7] Wei CHEN, Gregory AIST, Jack MOSTOW. Generating Questions Automatically from Informational Text.
- [8] Wenxing Hong, Siting Zheng, Huan Wang. Dynamic User Profile-Based Job Recommender. *The 8th International Conference on Computer Science & Education (ICCSE 2013)*, 2013.
- [9] Xuchen Yao, Gosse Bouma, Yi Zhang. Semantics-based Question Generation and Implementation. 2012.